

# NGS: Read Mapping to Reference Genome


**Daniel Sobral / José Ferrão**

# Learning Objectives

- Overall understanding of the data preprocessing steps
  - Alignment of NGS reads (fastq) to a reference sequence
    - The SAM / BAM format
    - Secondary and supplementary alignments
  - Marking Duplicates
  - Base Quality Recalibration
- Visualizing results and assessing quality


# Data pre-processing for variant discovery

## DNA Extraction




# Data pre-processing for variant discovery


Whole Genome



Targeted (eg. WES)



DNA Extraction



Eg. TruSight One Enrichment Panel

# How to calculate coverage

- WGS 30x coverage, 150bp read pairs
  - $30 \times 3 \times 10^9 / (150 \times 2) = 3 \times 10^8 = 300$  million read pairs
- WES (~1-2% genome) 30x coverage, 150bp read pairs
  - 6 million read pairs (theoretical minimum, but usually more)

# Genome (WGS) vs Exome (WES)

## Advantages of WES over WGS

- WES only focus on potentially actionable regions
  - WES is more economic and less compute intensive

## Disadvantages of WES

- WES can have primer-derived artifacts
  - eg. amplification bias, primer dimers, kit-dependent
- WES is very limited for structural variants

# Data pre-processing for variant discovery

Raw Unmapped Reads


uBAM or FASTQ



# Data pre-processing for variant discovery

Raw Unmapped Reads

uBAM or FASTQ






# Data pre-processing for variant discovery

Raw Unmapped Reads

uBAM or FASTQ

(fasta file)



Reference Genome



geneA AAGCGATGACTGCAT geneB GCACAGTTGTGT geneC GTTTTCACGTGAC

DNA Extraction

Sequencing



TGCTCAGTT



TGAC

ATGGAGT

GTTGT

("clean" fastQ)

QC and pre-processing

TGCTCAGTTA

TGAC

ATGGAGT

TTT

CGTTGT

("raw" fastQ)



# What genome reference to use?

- GRCh37 (hg19)
  - The “classical”, for which most resources were built (1000 genomes, etc...)
- **GRCh38 (hg38)**
  - **The recommended one. Most resources have been converted to this now**
- T2T-CHM13
  - Most complete genome, but will take time for resources to be available

Future: pangenome graphs

A lot of investment is going here, but several challenges remain

Mostly relevant for Structural Variants

# What genome reference to use?

Things to consider:

- Mask PAR regions in Y chromosome
- Include decoys (eg. EBV sequences)
- Include alternative sequences such as HLA loci
  - Though these may need to be analyzed specifically

<https://gatk.broadinstitute.org/hc/en-us/articles/360035890951-Human-genome-reference-builds-GRCh38-or-hg38-b37-hg19>

<https://gatk.broadinstitute.org/hc/en-us/articles/360035890951>


# Data pre-processing for variant discovery

Raw Unmapped Reads

uBAM or FASTQ

(fasta file)

Reference Genome




geneA geneB geneC  
AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

TGCTCAGTT

Alignment

DNA Extraction

Sequencing



TGCTCAGTT

TGAC

ATGGAGT

GTTGT

("clean" fastQ)

QC and pre-processing

TGCTCAGTTA

TGAC

ATGGAGT

TTT

CGTTGT

("raw" fastQ)


# Data pre-processing for variant discovery

Raw Unmapped Reads

uBAM or FASTQ

(fasta file)

Reference Genome



TGCTCAGTT

geneA

geneB

geneC


AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

Primary Alignment

Secondary Alignment

DNA Extraction

Sequencing



TGCTCAGTT

TGAC

ATGGAGT

GTTGT

("clean" fastQ)

QC and pre-processing

TGCTCAGTTA

TGAC

ATGGAGT

TTT

CGTTGT

("raw" fastQ)


# Data pre-processing for variant discovery

Raw Unmapped Reads

uBAM or FASTQ

(fasta file)

Reference Genome



ACTCCAT

GTTGT

TGCTCAGTT

geneA


geneB

geneC

AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

DNA Extraction

Sequencing



TGCTCAGTT

TGAC

ATGGAGT

GTTGT

("clean" fastQ)

QC and pre-processing

TGCTCAGTTA

TGAC

ATGGAGT

TTT

CGTTGT

("raw" fastQ)


# Data pre-processing for variant discovery

Raw Unmapped Reads

uBAM or FASTQ

(fasta file)

Reference Genome




TCCATGC AGTTGTGT  
 ACTCCAT GTTGT  
 AAGCGATG TGCTCAGTT GTGTGTTT-CA  
 geneA geneB geneC  
 AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

(SAM/BAM)

DNA Extraction

Sequencing



TGCTCAGTT TGAC ATGGAGT GTTGT  
 QC and pre-processing  
 (ACTCCAT)

("clean" fastQ)

TGCTCAGTTA TGAC ATGGAGT TTT CGTTGT ("raw" fastQ)



# SAM/BAM format

A file format to represent alignments

BAM -> binary form of SAM

```

Coor      12345678901234 5678901234567890123456789012345
ref       AGCATGTTAGATAA**GATAGCTGTGCTAGTAGGCAGTCAGCGCCAT

+r001/1      TTAGATAAAGGATA*CTG
+r002      aaaAGATAA*GGATA
+r003      gcctaAGCTAA
+r004      ATAGCT.....TCAGC
-r003      ttagctTAGGC
-r001/2      CAGCGGCAT
    
```



```

@HD VN:1.6 SO:coordinate
@SQ SN:ref LN:45
r001 99 ref 7 30 8M2I4M1D3M = 37 39 TTAGATAAAGGATACTG
r002 0 ref 9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA
r003 0 ref 9 30 5S6M * 0 0 GCCTAAGCTAA
r004 0 ref 16 30 6M14N5M * 0 0 ATAGCTTCAGC
r003 2064 ref 29 17 6H5M * 0 0 TAGGC
r001 147 ref 37 30 9M = 7 -39 CAGCGGCAT
    
```

<https://samtools.github.io/hts-specs/SAMv1.pdf>

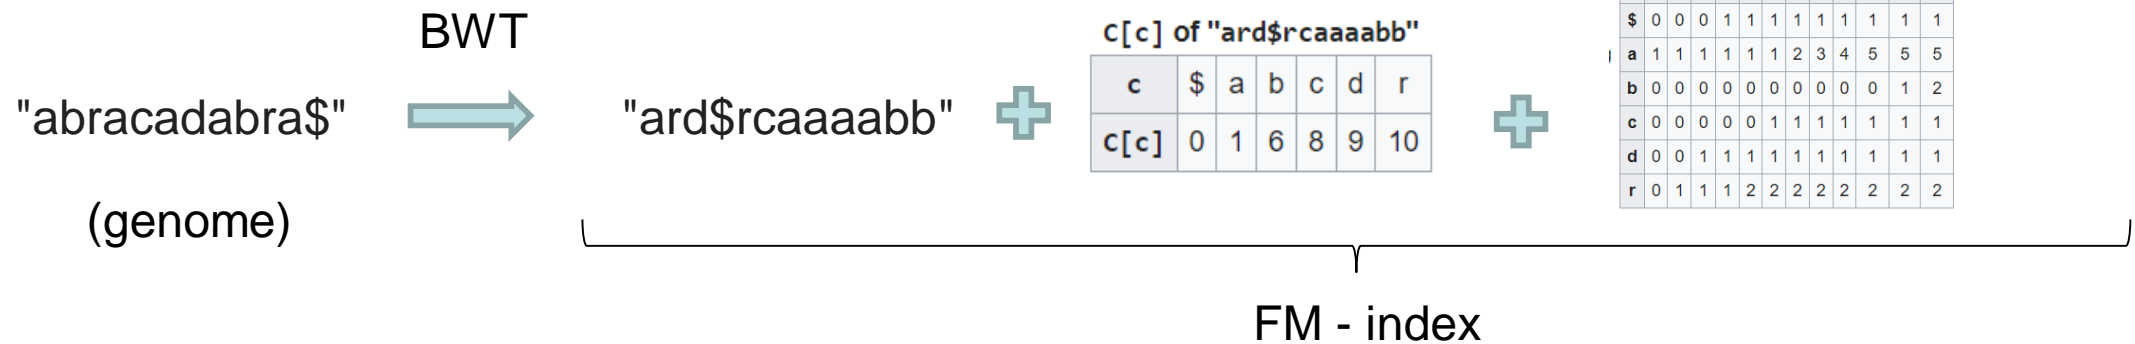


# How alignment is made in practice

BWA (Burrows-Wheeler Aligner) is the most popular tool for WGS/WES

- Align millions of short reads to a human-sized genome in minutes

It is based on the FM-index of the Burrows-Wheeler Transform



<https://en.wikipedia.org/wiki/FM-index>


# How alignment is made in practice

Genome BWT FM-index needs to be created (only once)

\$ bwa index genome.fasta

(genome.fasta file)

Reference Genome



BWT



c	\$	a	b	c	d	r
c[c]	0	1	6	8	9	10



	a	r	d	\$	r	c	a	a	a	a	b	b
	1	2	3	4	5	6	7	8	9	10	11	12
\$	0	0	0	1	1	1	1	1	1	1	1	1
a	1	1	1	1	1	1	2	3	4	5	5	5
b	0	0	0	0	0	0	0	0	0	1	2	
c	0	0	0	0	0	1	1	1	1	1	1	1
d	0	0	1	1	1	1	1	1	1	1	1	1
r	0	1	1	1	2	2	2	2	2	2	2	2

$3 \times 10^9$  bases

AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

FM - index

# How alignment is made in practice

Basic command to generate alignments with BWA:

```
$ bwa mem genome.fasta reads_R1.fastq(.gz) reads_R2.fastq(.gz)>output.sam
```

reads\_R1.fastq  
TGCTCAGTT

reads\_R2.fastq  
ACGTCCGA

BWT

c	\$	a	b	c	d	r
c[c]	0	1	6	8	9	10

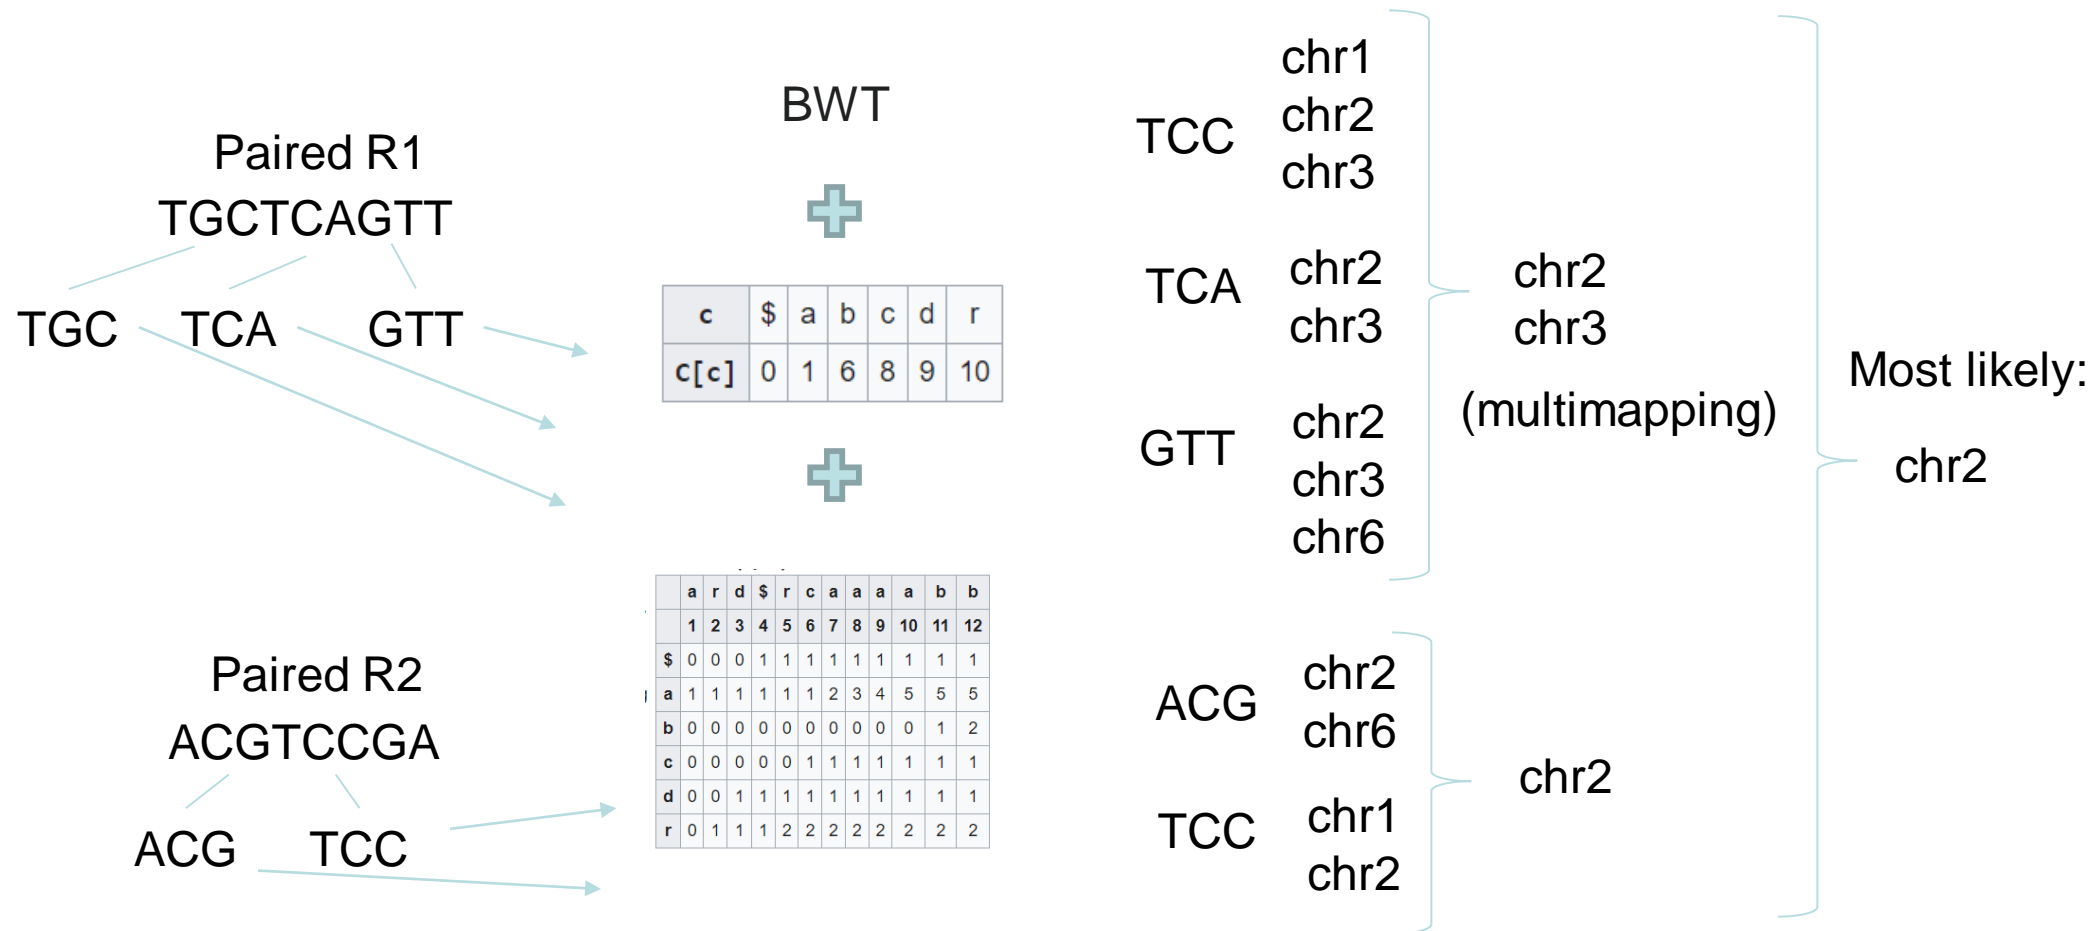
a	r	d	\$	r	c	a	a	a	a	b	b
1	2	3	4	5	6	7	8	9	10	11	12
\$	0	0	0	1	1	1	1	1	1	1	1
a	1	1	1	1	1	1	2	3	4	5	5
b	0	0	0	0	0	0	0	0	0	1	2
c	0	0	0	0	0	1	1	1	1	1	1
d	0	0	1	1	1	1	1	1	1	1	1
r	0	1	1	1	2	2	2	2	2	2	2

Output in SAM format

TGCTCAGTT Chr2 position xxx  
ACGTCCGA Chr2 position yyy

```
@HD VN:1.6 SO:coordinate
@SQ SN:ref LN:45
r001 99 ref 7 30 8M2I4M1D3M = 37 39 TTAGATAAAGGATACTG
r002 0 ref 9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA
r003 0 ref 9 30 5S6M * 0 0 GCCTAAGCTAA
r004 0 ref 16 30 6M14N5M * 0 0 ATAGCTTCAGC
r003 2064 ref 29 17 6H5M * 0 0 TAGGC
r001 147 ref 37 30 9M = 7 -39 CAGCGGCAT
```

# How alignment is made in practice





# How alignment is made in practice: Summary

- Special algorithms are used to have fast alignments
  - They are not guaranteed to be perfect but most of the time they are very good
- A read can map equally well to multiple regions (multimappings)
  - BWA reports **one primary alignment (randomly chosen) with mapping quality of 0**
  - Depending on the software, it can generate **secondary alignments**
  - Information of paired reads are used to disambiguate multimappings if possible
- Alignments are made piece-wise (a read is split in segments)
  - A read alignment can be split in a primary and **supplementary alignment(s)**
    - Eg. splicing in RNA-Seq; large deletions
  - Sometimes, only a part of the read is aligned (the rest is “masked”/hidden)
    - Particularly in repetitive areas this can lead to false alignments

# Practical Exercise

You will only need to run alignment using 'bwa mem', like this:

```
$ bwa mem -t 1 -M -R
```

```
'@RG\tID:sampleid\tPU:FlowCell.Lane.Sample\tSM:samplename\tPL:ILLUMINA\tLB:samplename\tCN:UTI' genoma.fasta reads_R1.fastq.gz reads_R2.fastq.gz > output.sam
```

-M: Mark split reads as secondary (relevant for detecting duplicates)

-R: Read Group (to associated specific metadata to each read)

ID: Identifier; PU: Platform Code (FlowCell.Lane.Sample); SM: Sample; PN: LB: Library; CN: Center Name (Sequencing Team) ; other fields possible (see SAM specification)

-t: number of "threads" (CPU processors) to use

# Practical Exercise

‘bwa mem’ generates a SAM file. To generate a more compact BAM file that is used to visualize and in variant calling, we need to process the SAM:

- Sort SAM by coordinates using ‘samtools sort’

```
$ samtools sort file.sam -o output_sorted.sam
```

- Convert SAM to BAM using ‘samtools view’

```
$ samtools view -Sb file.sam -o output.bam
```

-S: input file is SAM; -b: output file is bam

- Index the BAM so it can be used by other programs using ‘samtools index’

```
$ samtools sort file.bam -o output_sorted.bam
```

(note: this has nothing to do with ‘bwa index’)

# Practical Exercise

- There are three samples, use data from any/all:
  - Align selected reads against a specific chromosome
    - What is the identifier of the first aligned read? Where does it align?
  - Sort the aligned reads
    - What is the identifier of the first aligned read? Where does it align?
  - Convert to a bam and index the bam

**Commands to run are in the github**



# Interlude


# Data pre-processing for variant discovery

Raw Unmapped Reads

uBAM or FASTQ

(fasta file)

Reference Genome



TCCATGC AGTTGTGT  
 ACTCCAT GTTGT  
 AAGCGATG TGCTCAGTT GTGTGTTT-CA  
 geneA geneB geneC  
 AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

Alignments  
(SAM/BAM)

DNA Extraction

Sequencing



TGCTCAGTT TGAC ATGGAGT GTTGT  
 QC and pre-processing  
 (ACTCCAT)

("clean" fastQ)

TGCTCAGTTA TGAC ATGGAGT TTT CGTTGT ("raw" fastQ)

# Data pre-processing for variant discovery

## Duplicated Reads

- Duplicate reads (same fragment) can appear
  - In library preparation during amplification (eg. WES)
  - In the amplification process while sequencing (optical duplicates)
  - They can generate “false” mutations, or at a false frequency

```

AAGCGATG      AGTTGTGT
AAGCGATG      AGTTGTGT
      TCCATGC      GTTGT
      ACTCCAT
      TGCTCAGTT
      GTGTGTTT-CA

AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC
      geneA      geneB      geneC
    
```

# Data pre-processing for variant discovery

## Duplicated Reads

Read pairs need to be taken in consideration



# Data pre-processing for variant discovery

## Duplicated Reads

Duplicates are identified by position (alignment and/or flow-cell), not by sequence

```

AAGCTATG
AAGCGATG
AAGCGATG TCCATGC AGTTGTGT
          ACTCCAT  AGTTGTGT
                GTTGT
                TGCTCAGTT
                GTGTGTTT-CA

AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC
  geneA      geneB      geneC
  
```

# Data pre-processing for variant discovery

## Duplicated Reads: optical duplicates

@MN00723:33:000H3MCVT:1:11102:7591:1087 1:N:0:47

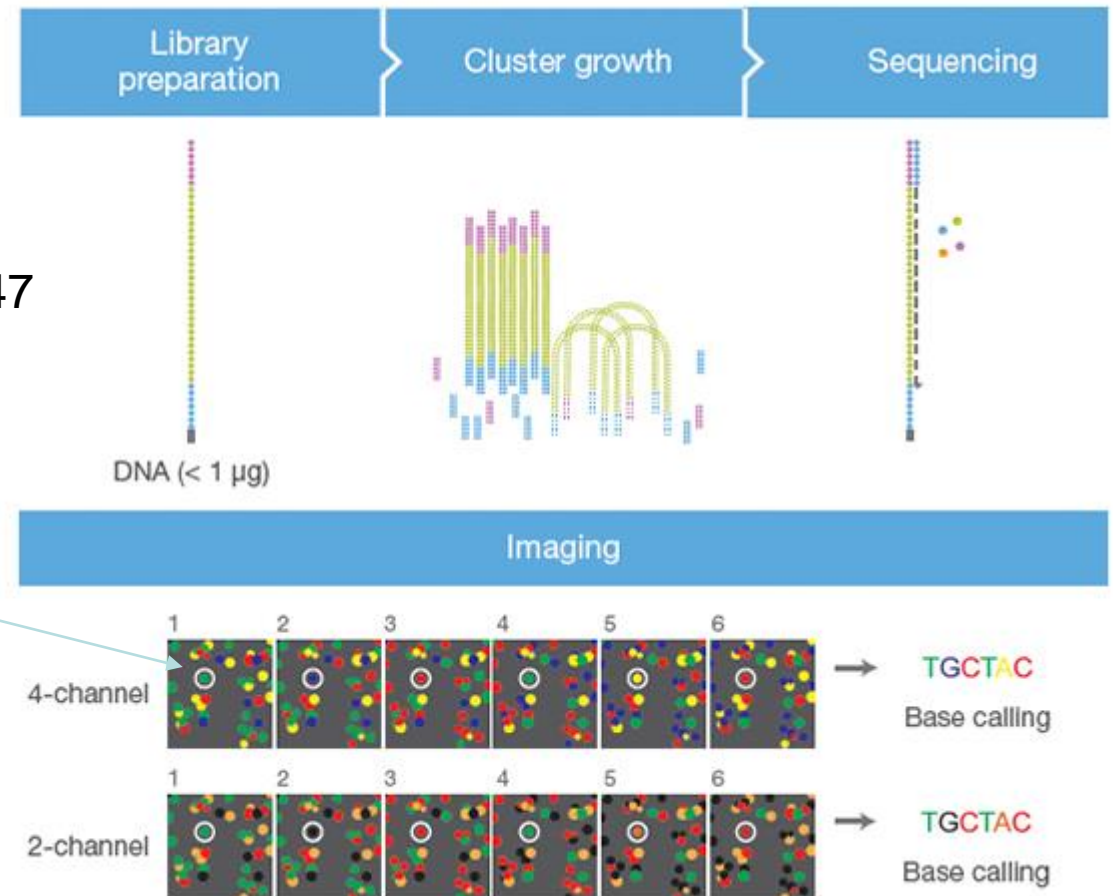
Machine

Flow Cell

Lane

Position

Optical duplicates are duplicate reads that are very close in the flow cell



# Data pre-processing for variant discovery

## Duplicated Reads

- **The recommended practice is to ignore duplicates**
  - Only consider one of the duplicates for variant calling
    - Usually the one with the best quality
  - This may remove good information (eg. with high coverage, targeted)
  - Duplicates are marked and later ignored (or not)
  - Benefits of marking duplicates not always obvious
    - <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC4965708/>
    - Eg. one can chose to only ignore reads marker as optical duplicates

# Data pre-processing for variant discovery

## Duplicated Reads

We will run this process using 'gatk MarkDuplicates'

It takes a BAM file and annotates duplicated reads:

```
> gatk MarkDuplicates -I=sample.bam -O=sample_marked.bam -M=sample_marked_metrics.txt
```

It generates another BAM file with the same number of alignments, but alignments determined as duplicated have an annotation that can be used.

It generates a file with metrics from the analysis (described in next slide)



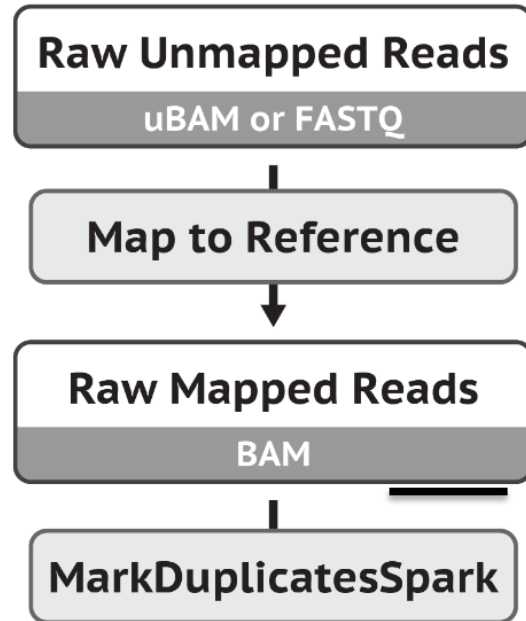
# Data pre-processing for variant discovery

## Duplicated Reads

### Metrics

Field	Description
LIBRARY	The library on which the duplicate marking was performed.
UNPAIRED_READS_EXAMINED	The number of mapped reads examined which did not have a mapped mate pair, either because the read is unpaired, or the read is paired to an unmapped mate.
READ_PAIRS_EXAMINED	The number of mapped read pairs examined. (Primary, non-supplemental)
SECONDARY_OR_SUPPLEMENTARY_RDS	The number of reads that were either secondary or supplementary
UNMAPPED_READS	The total number of unmapped reads examined. (Primary, non-supplemental)
UNPAIRED_READ_DUPLICATES	The number of fragments that were marked as duplicates.
READ_PAIR_DUPLICATES	The number of read pairs that were marked as duplicates.
READ_PAIR_OPTICAL_DUPLICATES	The number of read pairs duplicates that were caused by optical duplication. Value is always < READ_PAIR_DUPLICATES, which counts all duplicates regardless of source.
PERCENT_DUPLICATION	The fraction of mapped sequence that is marked as duplicate.
ESTIMATED_LIBRARY_SIZE	The estimated number of unique molecules in the library based on PE duplication.

# Data pre-processing for variant discovery



Now we have a BAM with potential duplicated alignments marked

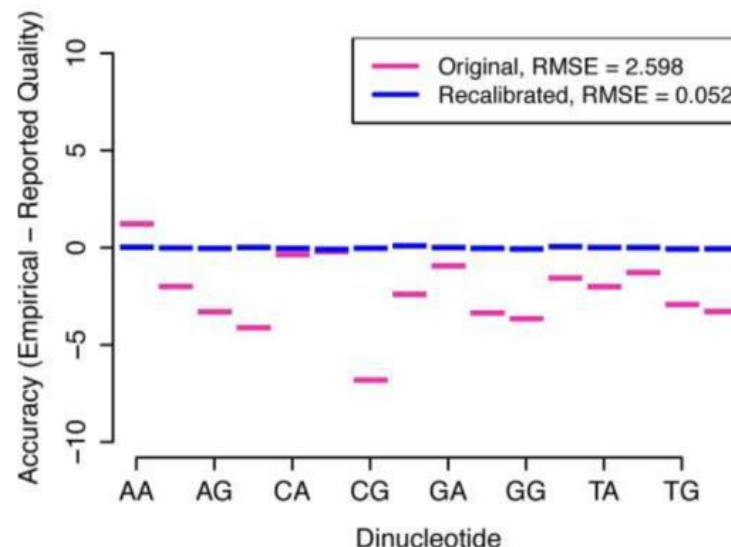
# Data pre-processing for variant discovery

## Base Quality Recalibration

Base Quality not very precise

Depends on several factors:

- Sample Quality (DNA)
- Nucleotide context
- Machine and cycle of sequencing
- Type of variant (SNP or Indel)



The machine oftens does not estimate correctly the base quality score

# Data pre-processing for variant discovery

## Base Quality Recalibration

Use list of known variants to estimate correct quality values

- All bases different from reference **not in the provided list** of known variants are considered to be errors



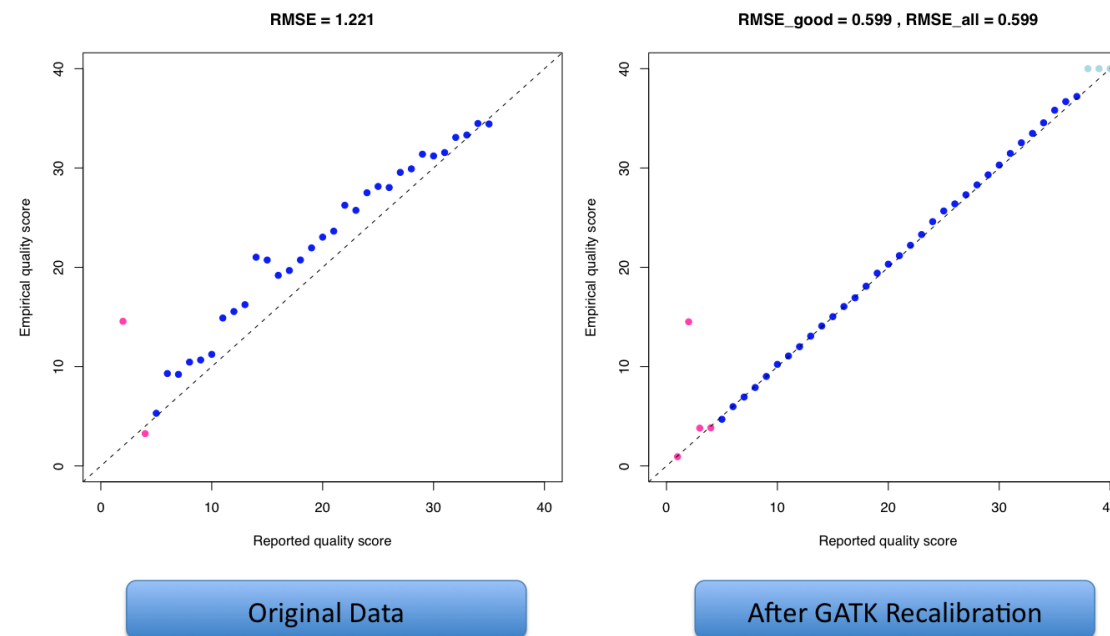
# Data pre-processing for variant discovery

## Base Quality Recalibration

### Base Quality Recalibration:

The covariates being used here:

- ReadGroupCovariate
- QualityScoreCovariate
- ContextCovariate
- CycleCovariate



<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3083463/>

<https://www.youtube.com/watch?v=L4D1dwES9s8>

<https://gatk.broadinstitute.org/hc/en-us/articles/360035890531-Base-Quality-Score-Recalibration-BQSR->

# Data pre-processing for variant discovery

## Base Quality Recalibration

We will run this process using 'gatk BaseRecalibrator / ApplyBQSR'

First, we need to generate a table with estimated base qualities split by covariate

```
> gatk BaseRecalibrator -I sample_marked.bam -R genome.fasta --known-sites known_snps.vcf
--known-sites known_indels.vcf [ --intervals target_positions.bed --interval-padding 100 ] -O
sample_marked_baserecalibrator_report.txt
```

In the case of WES we need to pass the relevant regions (--intervals and --interval-padding)

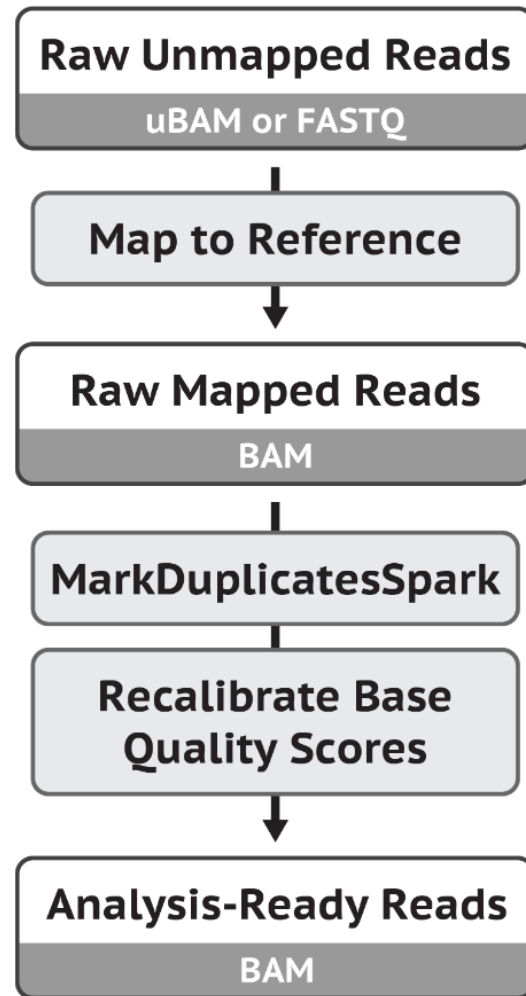
Next, we need to apply the inferred re-estimated qualities to the BAM file:

```
>gatk ApplyBQSR -I sample_marked.bam -R genome.fasta --bqsr-recal-file
sample_marked_baserecalibrator_report.txt -O sample_marked_baserecalibrator.bam
```

This will generate a final BAM, with marked duplicates and calibrated base qualities

# Data pre-processing for variant discovery

GATK  
Best  
Practices



These steps can  
be used with other  
variant callers

<https://gatk.broadinstitute.org/hc/en-us/articles/360035535912-Data-pre-processing-for-variant-discovery>

# Practical Exercise


- There are three samples, use data from any/all:
  - Mark duplicates using “gatk MarkDuplicates”
    - Describe the contents of the Metrics File
  - Apply Base Recalibration: gatk BaseRecalibrator / ApplyBQSR

**Commands to run are in the github**





# Interlude



# NGS: Quality Assessment and Visualization of Read Mappings

# Quality Assessment of Read Mappings

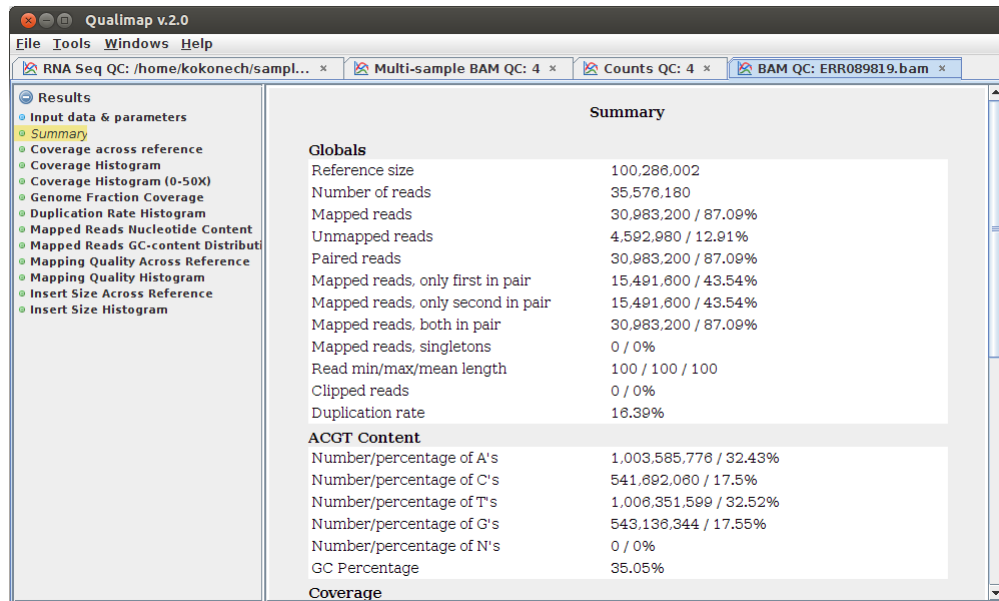
## Measures to consider:

- Mapped Reads / Unmapped Reads
  - Usually >80% for WGS; >90% for WES
- Total Mapped Reads / Coverage
  - Presence of Duplicates (depends on coverage):
    - Note that this is different from the duplicates inferred by FastQC
- Homogeneity in the coverage
- Distribution of Fragment Size (paired-end)
- Coverage of Targeted Regions (in case of eg. WES)

# Quality Assessment of Read Mappings

- Qualimap: similar to FastQC, but for BAM files

<http://qualimap.conesalab.org/>



# Practical example

Qualimap is a tool that can easily be run in Windows:

We will use 'qualimap bamqc':

```
>qualimap bamqc --bam sample_marked_recalibrated.bam --outdir . --outfile  
sample_marked_recalibrated_qualimap.pdf --outformat PDF [ -gff  
targets_position.bed ]
```

In the case of WES you need to pass a file with relevant intervals

It will generate a PDF (or HTML) with results, like FASTQC

# Practical Exercise

Run qualimap bamqc for any/all of the samples

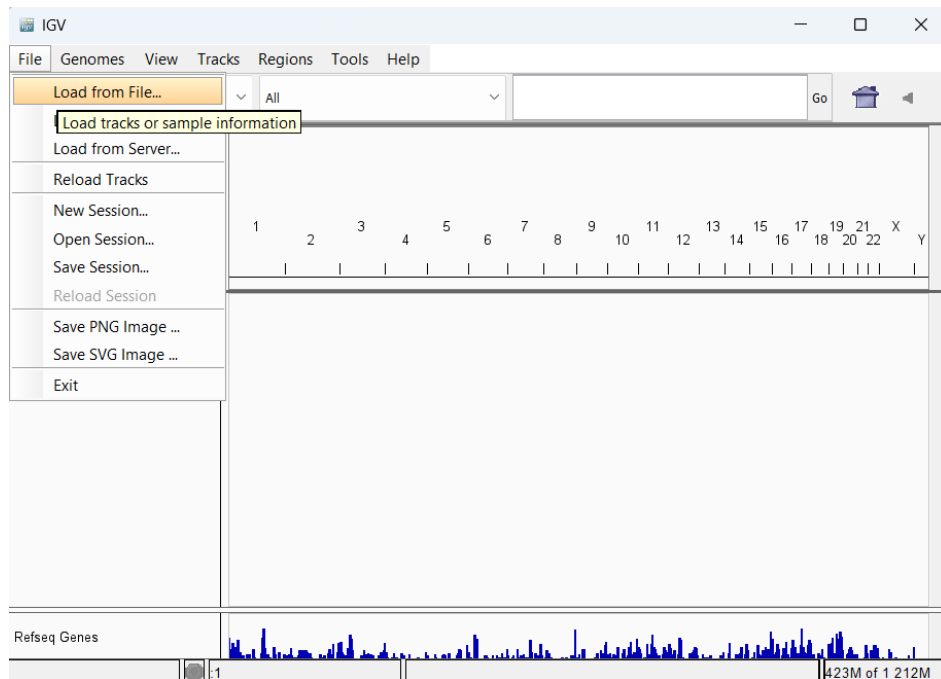
- Compare with the metrics from GATK markduplicates
- What is the distribution of fragment length?
- If you have time, compare reports of BAMs from each of the steps

**Commands to run are in the github**

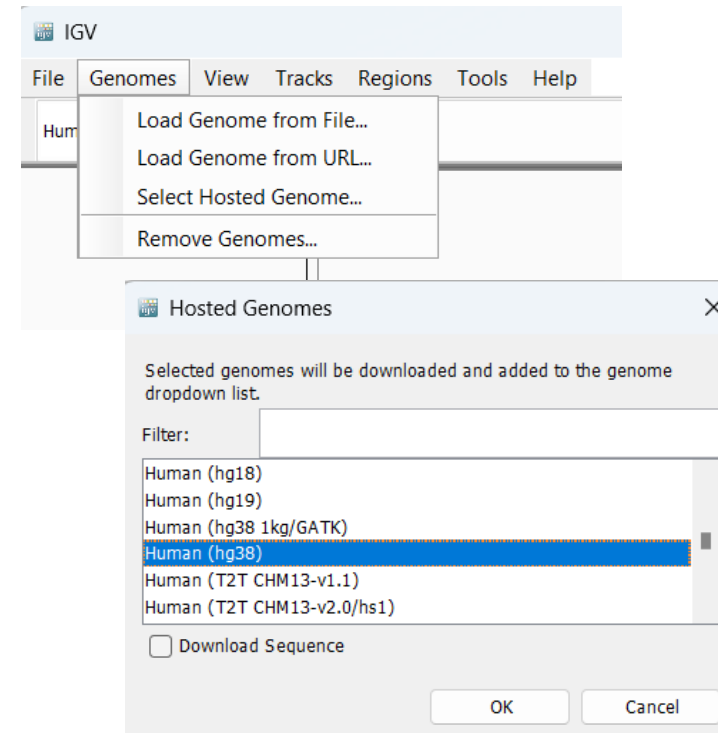
# Visualization of Read Mappings with IGV

Load Data in IGV:


- It can be BAM, but also VCF



If genome is not loaded:



# Visualization of Read Mappings with IGV



Coverage


Alignments

Annotation  
(genes)



# Visualization of Read Mappings with IGV


## Example mutations



## Example deletion



- IGV provides colors to signal unusual situations
  - Besides mutations, information from paired-end is also there



Pairs on different chromosomes

Insert Size Lengths

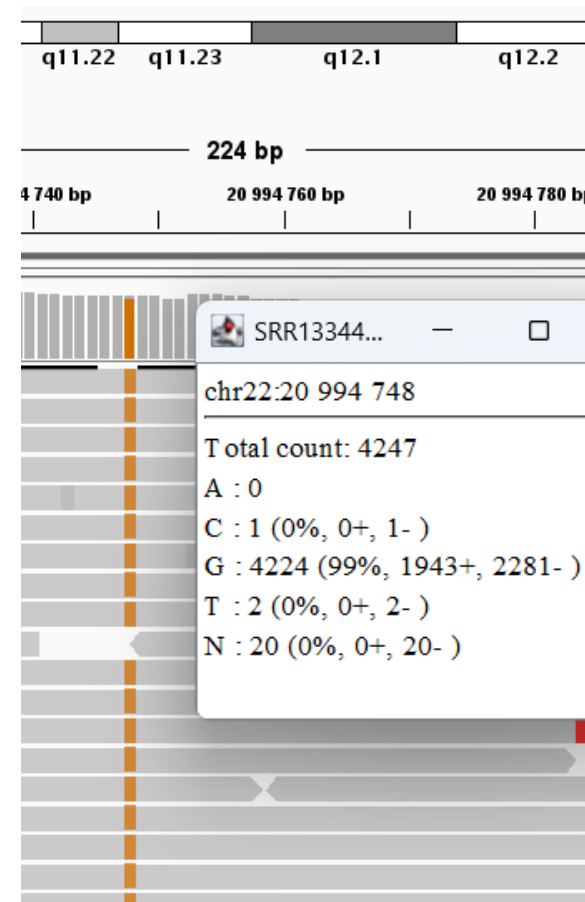
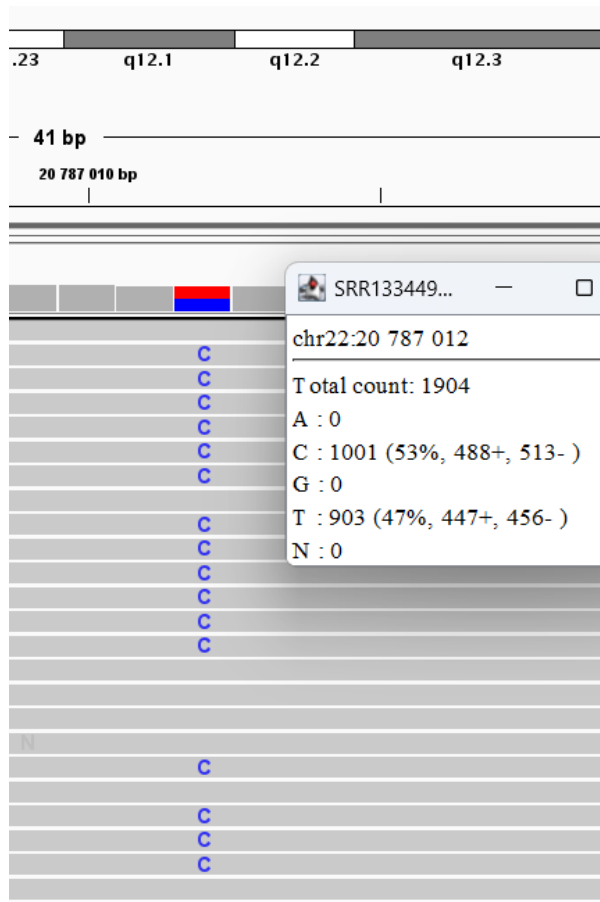
Category	Illumina		
LR			Normal
LL			Unusual
RR			Unusual
RL			Unusual

Pair orientation

<https://software.broadinstitute.org/software/igv/>

# Visualization of Read Mappings

- Example of mutations, showing evidence



# Practical Exercise

- Look at the alignments in IGV in specific positions:
  - Sample 1: 2:171868791
  - Sample 2: 6:121216049
  - Sample 3: X:129746638

Gather evidence that you think is relevant for the mutation:

- Number of reads, strand of reads, quality etc...