

NGS: Read Mapping to Reference Genome


Daniel Sobral / José Ferrão

Learning Objectives

- Overall understanding of the data preprocessing steps
 - Alignment of NGS reads (fastq) to a reference sequence
 - The SAM / BAM format
 - Marking Duplicates
 - Base Quality Recalibration
- Visualizing results and assessing quality


Data pre-processing for variant discovery

DNA Extraction




Data pre-processing for variant discovery


Whole Genome



Targeted (eg. WES)



DNA Extraction



Eg. TruSight One Enrichment Panel

How to calculate coverage

- WGS 30x coverage, 150bp read pairs
 - $30 \times 3 \times 10^9 / (150 \times 2) = 3 \times 10^8 = 300$ million read pairs
- WES (~1-2% genome) 30x coverage, 150bp read pairs
 - 6 million read pairs (theoretical minimum, but usually more)

Data pre-processing for variant discovery

Raw Unmapped Reads

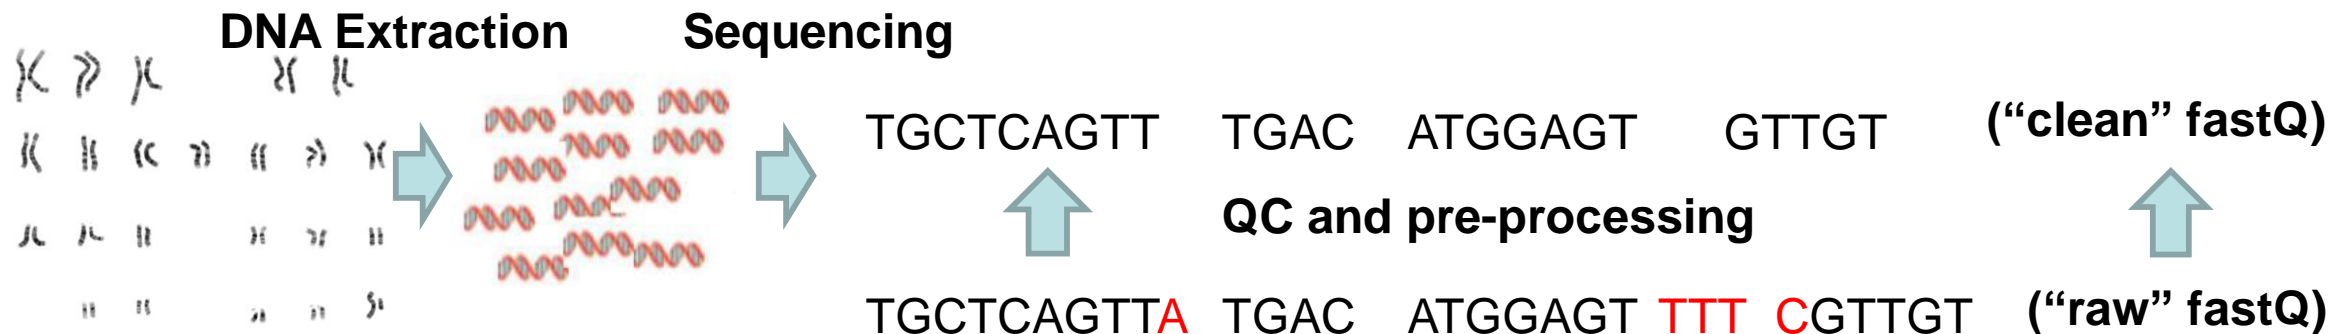
uBAM or FASTQ



Data pre-processing for variant discovery

Raw Unmapped Reads

uBAM or FASTQ




Data pre-processing for variant discovery

Raw Unmapped Reads

uBAM or FASTQ

(fasta file)


Reference Genome



geneA AAGCGATGACTGCAT geneB GCACAGTTGTGT geneC GTTTTCACGTGAC

DNA Extraction

Sequencing



TGCTCAGTT



TGAC

ATGGAGT

GTTGT

("clean" fastQ)

QC and pre-processing

TGCTCAGTTA

TGAC

ATGGAGT

TTT

CGTTGT

("raw" fastQ)



What genome reference to use?

- GRCh37 (hg19)
 - The “classical”, for which most resources were built (1000 genomes, etc...)
- **GRCh38 (hg38)**
 - **The recommended one. Most resources have been converted to this now**
- T2T-CHM13
 - Most complete genome, but will take time for resources to be available

Future: pangenome graphs

A lot of investment is going here, but several challenges remain

Mostly relevant for Structural Variants

What genome reference to use?

Things to consider:

- Mask PAR regions in Y chromosome
- Include decoys (eg. EBV sequences)
- Include alternative sequences such as HLA loci
 - Though these may need to be analyzed specifically

<https://gatk.broadinstitute.org/hc/en-us/articles/360035890951-Human-genome-reference-builds-GRCh38-or-hg38-b37-hg19>

<https://gatk.broadinstitute.org/hc/en-us/articles/360035890951>


Data pre-processing for variant discovery

Raw Unmapped Reads

uBAM or FASTQ

(fasta file)

Reference Genome




geneA geneB geneC

AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

TGCTCAGTT

DNA Extraction

Sequencing



TGCTCAGTT

TGAC

ATGGAGT

GTTGT

("clean" fastQ)

QC and pre-processing

TGCTCAGTTA

TGAC

ATGGAGT

TTT

CGTTGT

("raw" fastQ)


Data pre-processing for variant discovery

Raw Unmapped Reads

uBAM or FASTQ

(fasta file)

Reference Genome



TGCTCAGTT

geneA

geneB


geneC

AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

DNA Extraction

Sequencing

Alignment



TGCTCAGTT

TGAC

ATGGAGT

GTTGT

("clean" fastQ)

QC and pre-processing

TGCTCAGTTA

TGAC

ATGGAGT

TTT

CGTTGT

("raw" fastQ)


Data pre-processing for variant discovery

Raw Unmapped Reads

uBAM or FASTQ


(fasta file)

Reference Genome



ACTCCAT GTTGT
TGCTCAGTT

geneA geneB geneC
AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC




Data pre-processing for variant discovery

Raw Unmapped Reads

uBAM or FASTQ

(fasta file)

Reference Genome



TCCATGC AGTTGTGT
 ACTCCAT GTTGT
 AAGCGATG TGCTCAGTT GTGTGTTT-CA
 geneA geneB geneC
 AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

(SAM/BAM)


DNA Extraction

Sequencing

Alignment

("clean" fastQ)

("raw" fastQ)



TGCTCAGTT TGAC ATGGAGT GTTGT
 QC and pre-processing
 TGCTCAGTTA TGAC ATGGAGT TTT CGTTGT

(ACTCCAT)

SAM/BAM format

A file format to represent alignments

BAM -> binary form of SAM

```

Coor      12345678901234  5678901234567890123456789012345
ref       AGCATGTTAGATAA**GATAGCTGTGCTAGTAGGCAGTCAGCGCCAT

+r001/1      TTAGATAAAGGATA*CTG
+r002      aaaAGATAA*GGATA
+r003      gcctaAGCTAA
+r004      ATAGCT.....TCAGC
-r003      ttagctTAGGC
-r001/2      CAGCGGCAT
    
```



```

@HD VN:1.6 SO:coordinate
@SQ SN:ref LN:45
r001  99 ref  7 30 8M2I4M1D3M = 37 39 TTAGATAAAGGATACTG
r002   0 ref  9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA
r003   0 ref  9 30 5S6M      * 0 0 GCCTAAGCTAA
r004   0 ref 16 30 6M14N5M   * 0 0 ATAGCTTCAGC
r003 2064 ref 29 17 6H5M     * 0 0 TAGGC
r001  147 ref 37 30 9M      = 7 -39 CAGCGGCAT
    
```

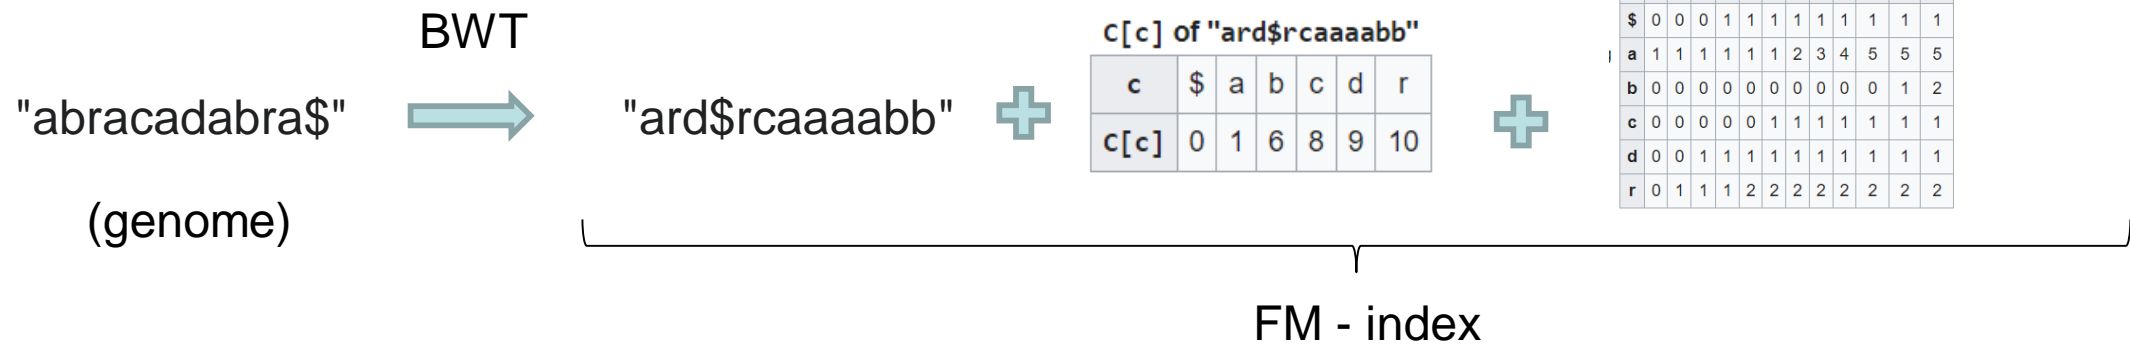
<https://samtools.github.io/hts-specs/SAMv1.pdf>

How alignment is made in practice

BWA (Burrows-Wheeler Aligner) is the most popular tool for WGS/WES

- Align millions of short reads to a human-sized genome in minutes

It is based on the FM-index of the Burrows-Wheeler Transform



<https://en.wikipedia.org/wiki/FM-index>


How alignment is made in practice

Genome BWT FM-index needs to be created (only once)

\$ bwa index genome.fasta

(genome.fasta file)

Reference Genome



3×10^9 bases



BWT



| c | \$ | a | b | c | d | r |
|------|----|---|---|---|---|----|
| c[c] | 0 | 1 | 6 | 8 | 9 | 10 |



| | a | r | d | \$ | r | c | a | a | a | a | b | b |
|----|---|---|---|----|---|---|---|---|---|----|----|----|
| | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 |
| \$ | 0 | 0 | 0 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| a | 1 | 1 | 1 | 1 | 1 | 1 | 2 | 3 | 4 | 5 | 5 | 5 |
| b | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 2 |
| c | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| d | 0 | 0 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| r | 0 | 1 | 1 | 1 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 |

FM - index

How alignment is made in practice

Command to generate alignments with BWA:

```
$ bwa mem genome.fasta reads_R1.fastq(.gz) reads_R2.fastq(.gz)
(BWT FM-index for genome.fasta needs to be created before)
```

reads_R1.fastq

TGCTCAGTT

reads_R2.fastq

ACGTCCGA

BWT

| c | \$ | a | b | c | d | r |
|------|----|---|---|---|---|----|
| c[c] | 0 | 1 | 6 | 8 | 9 | 10 |

| a | r | d | \$ | r | c | a | a | a | a | b | b |
|----|---|---|----|---|---|---|---|---|----|----|----|
| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 |
| \$ | 0 | 0 | 0 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| a | 1 | 1 | 1 | 1 | 1 | 1 | 2 | 3 | 4 | 5 | 5 |
| b | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 2 |
| c | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 1 | 1 | 1 | 1 |
| d | 0 | 0 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| r | 0 | 1 | 1 | 1 | 2 | 2 | 2 | 2 | 2 | 2 | 2 |

TGCTCAGTT

Chr2 position xxx

ACGTCCGA

Chr2 position yyy

How alignment is made in practice

Paired R1

TGCTCAGTT

TGC

TCA

GTT

BWT



| c | \$ | a | b | c | d | r |
|------|----|---|---|---|---|----|
| c[c] | 0 | 1 | 6 | 8 | 9 | 10 |



| | a | r | d | \$ | r | c | a | a | a | a | b | b |
|----|---|---|---|----|---|---|---|---|---|----|----|----|
| | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 |
| \$ | 0 | 0 | 0 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| a | 1 | 1 | 1 | 1 | 1 | 1 | 2 | 3 | 4 | 5 | 5 | 5 |
| b | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 2 |
| c | 0 | 0 | 0 | 0 | 0 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| d | 0 | 0 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| r | 0 | 1 | 1 | 1 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 |

Paired R2

ACGTCCGA

ACG

TCC

TCC

chr1
chr2
chr3

TCA

chr2
chr3

GTT

chr2
chr3
chr6

ACG

chr2
chr6

TCC

chr1
chr2

chr2
chr3

(multimapping)

chr2

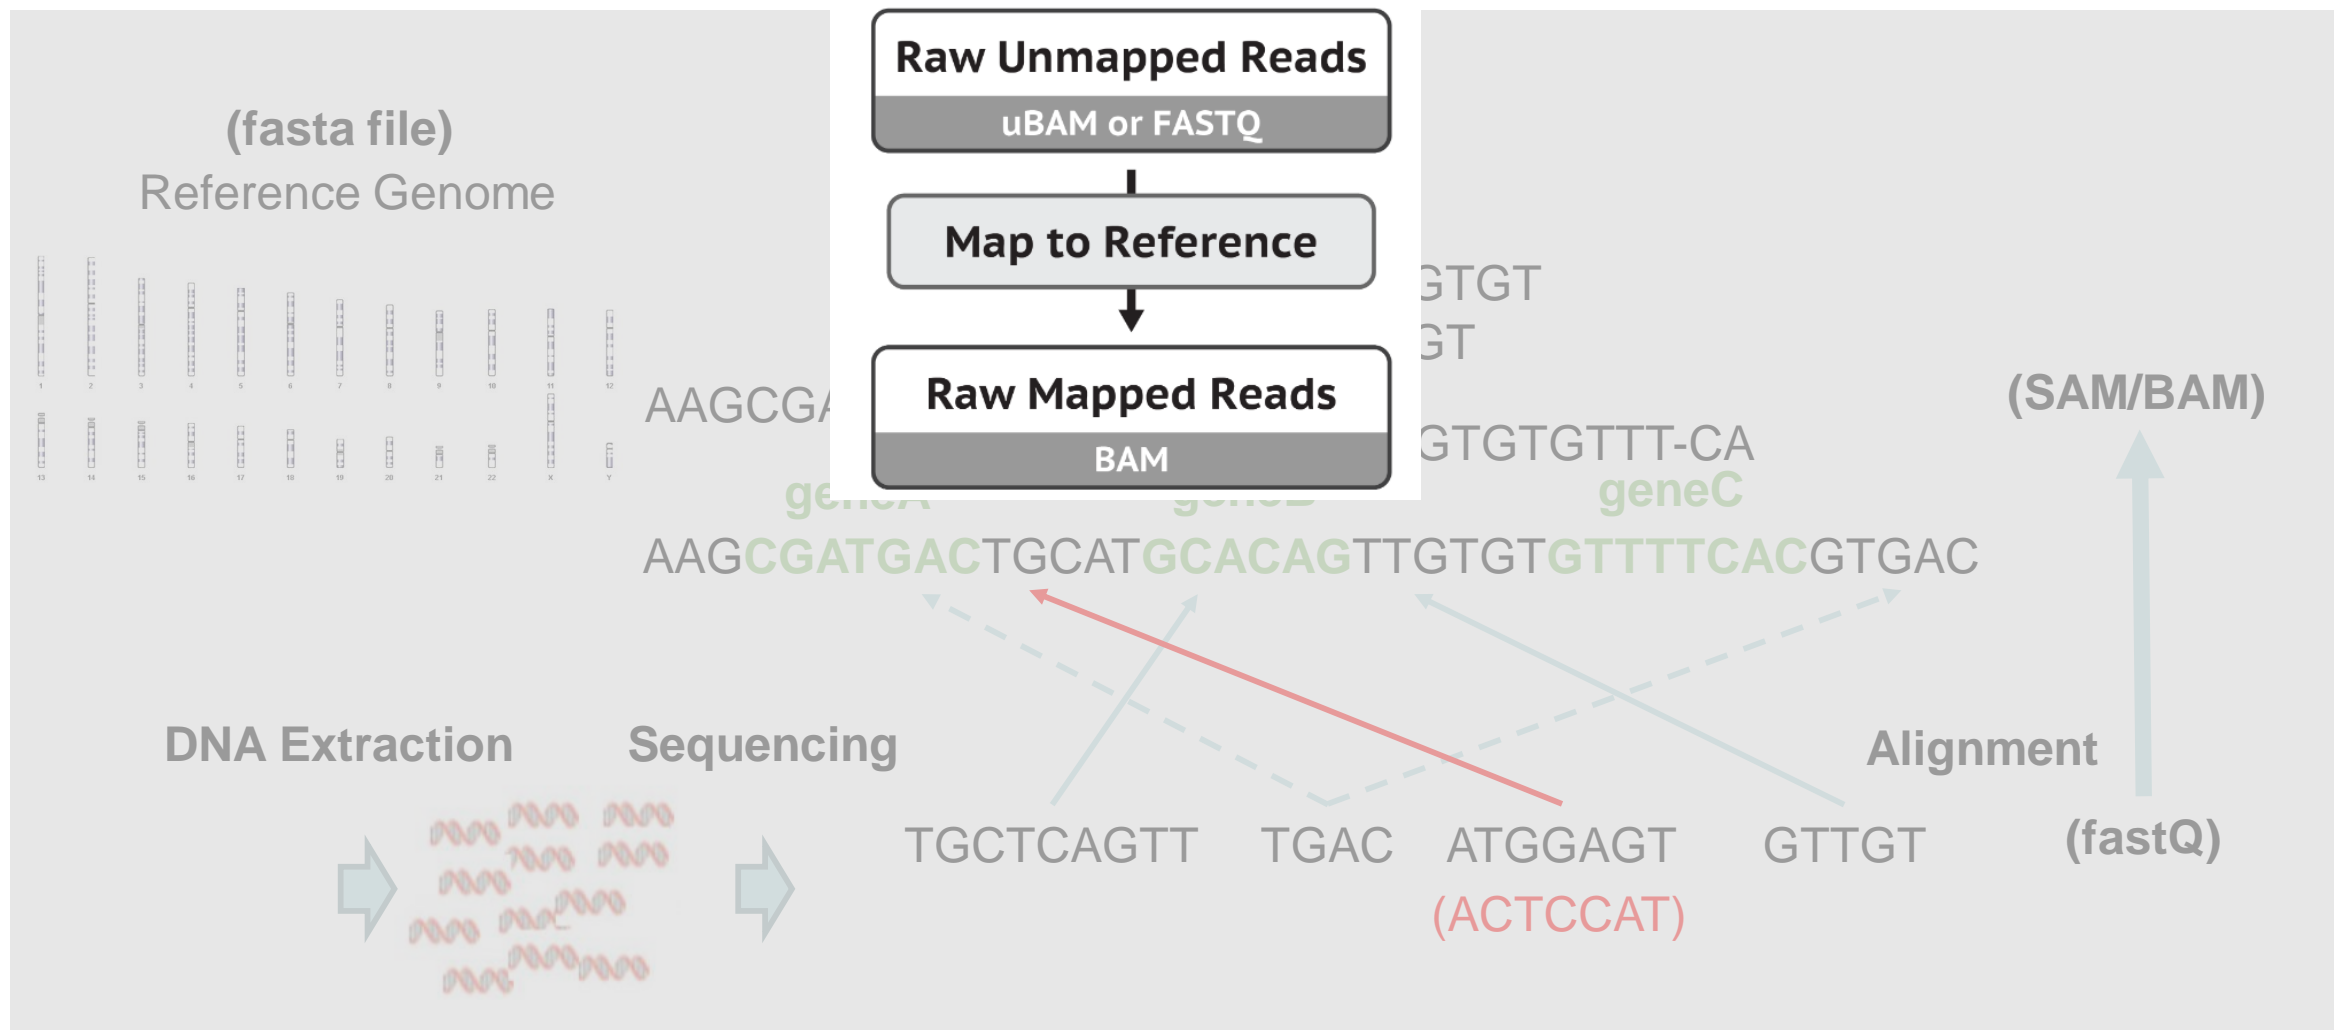
Most likely:
chr2



How alignment is made in practice

- Alignments are fast but may not be always perfect
 - Nonetheless it works very well, particularly for human
- Equally good mappings (multimappings) are marked as such
 - Mapping quality is set as 0 by default
- Alignments are made piece-wise
 - A read can be split in a principal and secondary alignment(s)
 - Eg. splicing in RNA-Seq; large insertions
 - Sometimes, only a part of the read is aligned (the rest is masked)
 - Particularly in repetitive areas this can lead to false alignments

Data pre-processing for variant discovery



Practical Example

GATK (Broad Institute) provides resources for their best practices protocol:

<https://console.cloud.google.com/storage/browser/genomics-public-data/resources/broad/hg38/v0/>

(already includes pre-built BWT FM-index)

```
> bwa index ficheiro_genoma.fasta
```

```
> bwa mem -t 1 -R
```

```
'@RG\tID:sampleid\tPU:FlowCell.Lane.Sample\tSM:samplename\tPL:ILLUMINA\tLB:samplename\tCN:UTI' -M ficheiro_genoma.fasta amostrareads_R1.fastq.gz amostrareads_R2.fastq.gz > amostra.sam
```

ID: Identifier; PU: Platform Code (FlowCell.Lane.Sample); SM: Sample; PN: LB: Library; CN: Center Name (Sequencing Team) ; other fields possible (see SAM specification)

Exercise

- Each group should have one sample (1,2,3)
 - How many reads do you have in your sample?
 - Align selected reads against a specific chromosome
 - What is the identifier of the first aligned read? Where does it align?
 - How many entries are in the SAM file? Is it the number of reads?
 - Can you see in the header information about the program being run
 - Can you see the read group information you added?
 - Sort the aligned reads
 - What is the identifier of the first aligned read? Where does it align?
 - Convert sam to a bam and index the bam
 - Compare the size of the files. What other files were generated?


Data pre-processing for variant discovery

Raw Unmapped Reads

uBAM or FASTQ

(fasta file)

Reference Genome



TCCATGC AGTTGTGT
 ACTCCAT GTTGT
 AAGCGATG TGCTCAGTT GTGTGTTT-CA
 geneA geneB geneC
 AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC

(SAM/BAM)


DNA Extraction

Sequencing

Alignment

("clean" fastQ)

("raw" fastQ)



TGCTCAGTT TGAC ATGGAGT GTTGT
 QC and pre-processing
 TGCTCAGTTA TGAC ATGGAGT TTT CGTTGT

(ACTCCAT)

Data pre-processing for variant discovery

Duplicated Reads

- Duplicate reads (same fragment) can appear
 - In library preparation during amplification (eg. WES)
 - In the amplification process while sequencing (optical duplicates)

```

AAGCGATG      AGTTGTGT
AAGCGATG      AGTTGTGT
      TCCATGC      GTTGT
      ACTCCAT
      TGCTCAGTT
      GTGTGTTT-CA

AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC
  geneA          geneB          geneC
    
```

Data pre-processing for variant discovery

Duplicated Reads

Read pairs need to be taken in consideration



Data pre-processing for variant discovery

Duplicated Reads

Duplicates are identified by position (alignment and/or flow-cell), not by sequence

```

AAGCTATG
AAGCGATG
AAGCGATG TCCATGC AGTTGTGT
          ACTCCAT  AGTTGTGT
                GTTGT
                TGCTCAGTT
                GTGTGTTT-CA

AAGCGATGACTGCATGCACAGTTGTGTGTTTTCACGTGAC
  geneA      geneB      geneC
    
```

Data pre-processing for variant discovery

Duplicated Reads: optical duplicates

@MN00723:33:000H3MCVT:1:11102:7591:1087 1:N:0:47


Machine

Flow Cell

Lane

Position

Optical duplicates are duplicate reads that are very close in the flow cell



Data pre-processing for variant discovery

Duplicated Reads

- **The recommended practice is to ignore duplicates**
 - Only consider one of the duplicates for variant calling
 - Usually the one with the best quality
 - This may remove good information (eg. with high coverage, targeted)
 - Duplicates are marked and later ignored (or not)
 - Benefits of marking duplicates not always obvious
 - <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC4965708/>
 - Eg. one can choose to only ignore reads marked as optical duplicates

Practical Exercise

```
> gatk MarkDuplicates -I=sample.bam -O=sample_marked.bam -M=sample_marked_metrics.txt
```

Alternative: samblaster

```
> samblaster sample.bam | samtools sort -o sample.marked.sorted.bam
```

Data pre-processing for variant discovery

Duplicated Reads

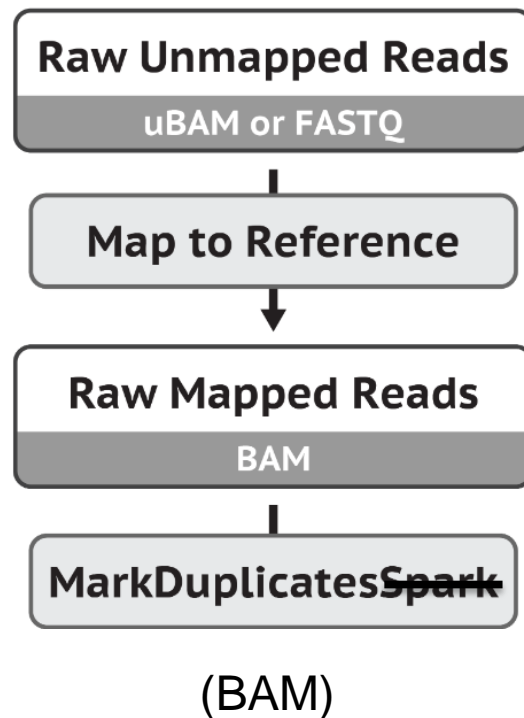
Metrics

| Field | Description |
|--------------------------------|--|
| LIBRARY | The library on which the duplicate marking was performed. |
| UNPAIRED_READS_EXAMINED | The number of mapped reads examined which did not have a mapped mate pair, either because the read is unpaired, or the read is paired to an unmapped mate. |
| READ_PAIRS_EXAMINED | The number of mapped read pairs examined. (Primary, non-supplemental) |
| SECONDARY_OR_SUPPLEMENTARY_RDS | The number of reads that were either secondary or supplementary |
| UNMAPPED_READS | The total number of unmapped reads examined. (Primary, non-supplemental) |
| UNPAIRED_READ_DUPLICATES | The number of fragments that were marked as duplicates. |
| READ_PAIR_DUPLICATES | The number of read pairs that were marked as duplicates. |
| READ_PAIR_OPTICAL_DUPLICATES | The number of read pairs duplicates that were caused by optical duplication. Value is always < READ_PAIR_DUPLICATES, which counts all duplicates regardless of source. |
| PERCENT_DUPLICATION | The fraction of mapped sequence that is marked as duplicate. |
| ESTIMATED_LIBRARY_SIZE | The estimated number of unique molecules in the library based on PE duplication. |

Exercise

- Run MarkDuplicates for your sample
 - There should be a metrics file generated and a new bam (we will discuss the content of the metrics file later)

Data pre-processing for variant discovery



Now we have alignments where each read may have an extra annotation if it is recognized as being duplicate.


Data pre-processing for variant discovery

Base Quality Recalibration

Base Quality not very precise

Depends on several factors:

- Sample Quality (DNA)
- Nucleotide context
- Machine and cycle of sequencing
- Type of variant (SNP or Indel)




The machine oftens does not estimate well the base quality score

Data pre-processing for variant discovery

Base Quality Recalibration

Use list of known variants to estimate correct quality values

- All bases different from reference **not in the provided list** of known variants are considered to be errors




Data pre-processing for variant discovery

Base Quality Recalibration

Base Quality Recalibration:

The covariates being used here:

- ReadGroupCovariate
- QualityScoreCovariate
- ContextCovariate
- CycleCovariate



<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3083463/>

<https://www.youtube.com/watch?v=L4D1dwES9s8>

<https://gatk.broadinstitute.org/hc/en-us/articles/360035890531-Base-Quality-Score-Recalibration-BQSR->

Data pre-processing for variant discovery

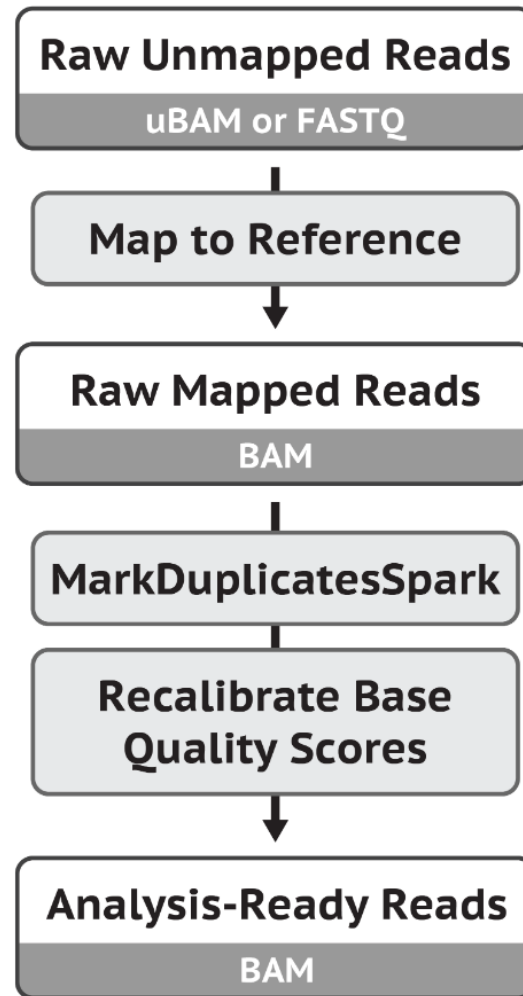
Base Quality Recalibration

```
> gatk BaseRecalibrator -I sample_marked.bam -R genome.fasta --known-sites known_snps.vcf  
--known-sites known_indels.vcf [ --intervals target_positions.bed --interval-padding 100 ] -O  
sample_marked_baserecalibrator_report.txt
```


```
> gatk ApplyBQSR -I sample_marked.bam -R genome.fasta --bqsr-recal-file  
sample_marked_baserecalibrator_report.txt -O sample_marked_baserecalibrator.bam
```

Data pre-processing for variant discovery

GATK
Best
Practices



These steps can
be used with other
variant callers



NGS: Quality Assessment and Visualization of Read Mappings

Quality Assessment of Read Mappings

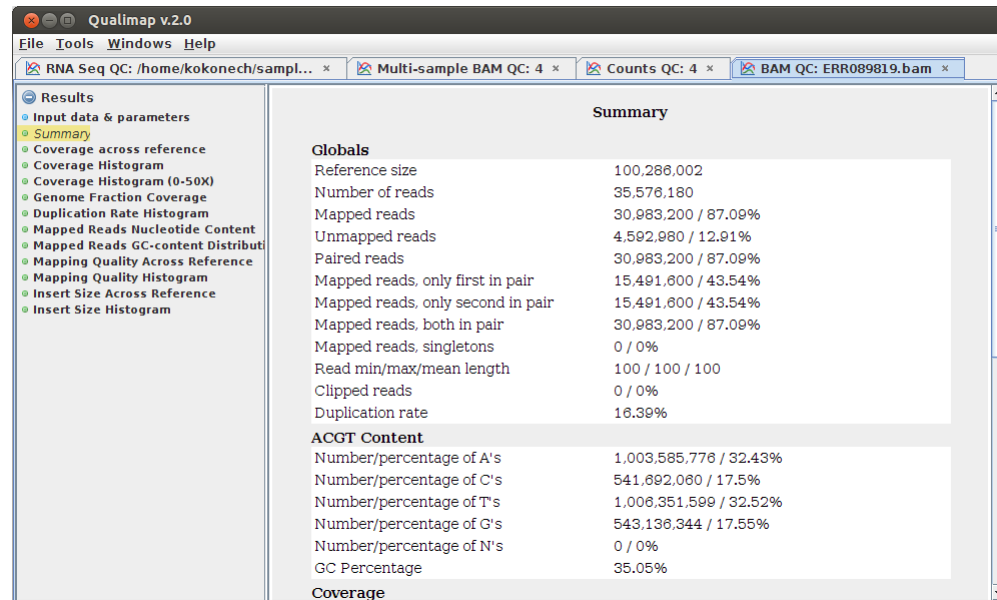
Measures to consider:

- Mapped Reads / Unmapped Reads
 - Usually >80% for WGS; >90% for WES
- Total Mapped Reads / Coverage
 - Presence of Duplicates (depends on coverage)
- Homogeneity in the coverage
- Distribution of Fragment Size (paired-end)
- Coverage of Targeted Regions (in case of eg. WES)

Quality Assessment of Read Mappings

- Qualimap

<http://qualimap.conesalab.org/>



Practical example

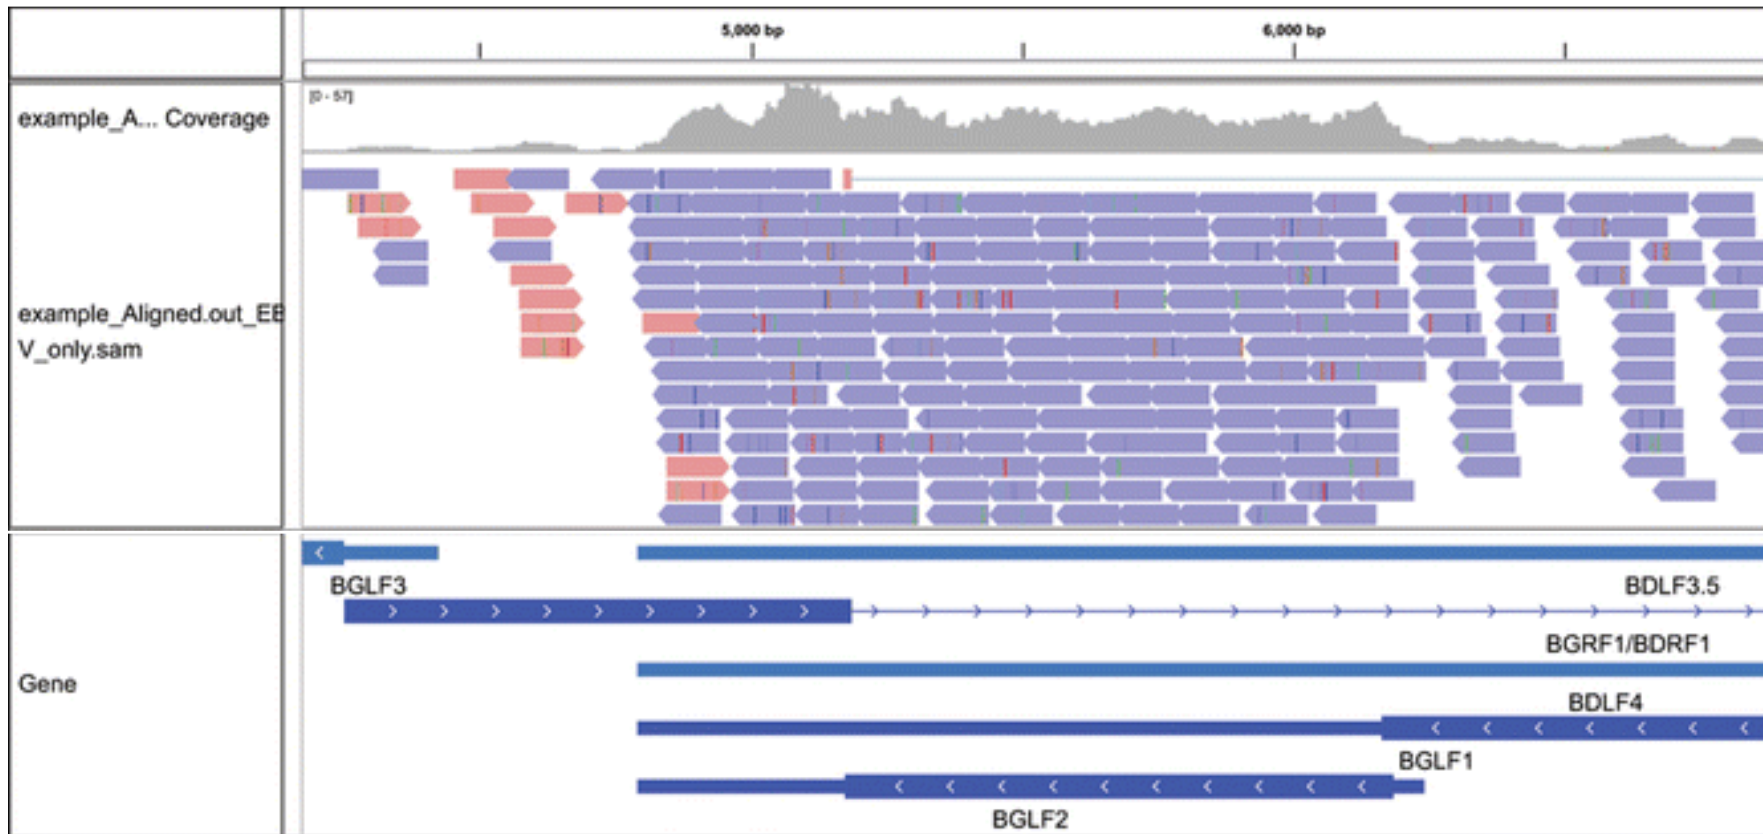
- Look at Metrics from GATK

```
>qualimap bamqc -bam sample_marked_recalibrated.bam -outdir . -outfile  
sample_marked_recalibrated_qualimap.pdf -outformat PDF [ -gff  
targets_position.bed ]
```

Exercise

- Look at Metrics from GATK for your sample
- Run qualimap bamqc for your sample and see output

Visualization of Read Mappings




Coverage

Alignments


Annotation
(genes)

Visualization of Read Mappings

Example mutations



- IGV provides colors to signal unusual situations
 - Besides mutations, information from paired-end is also there



Pairs on different chromosomes

Insert Size Lengths

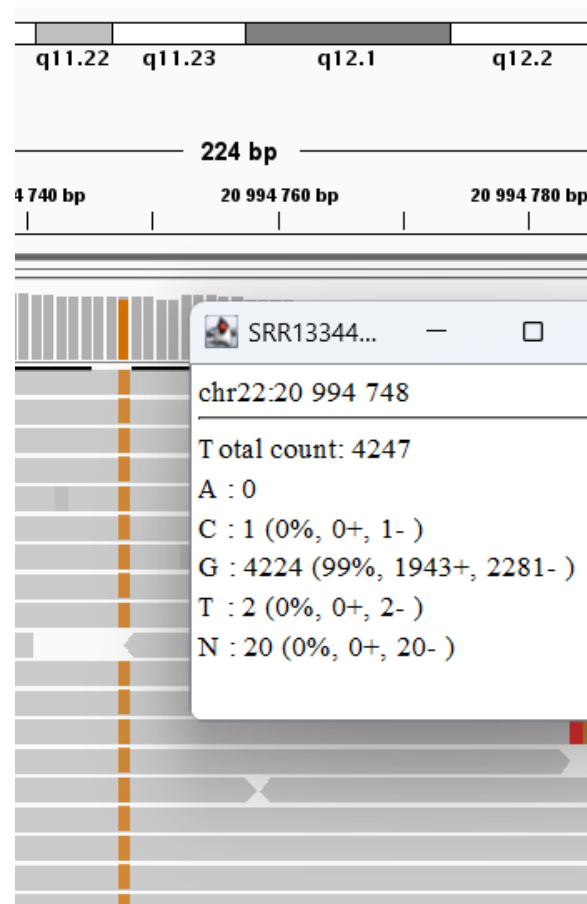
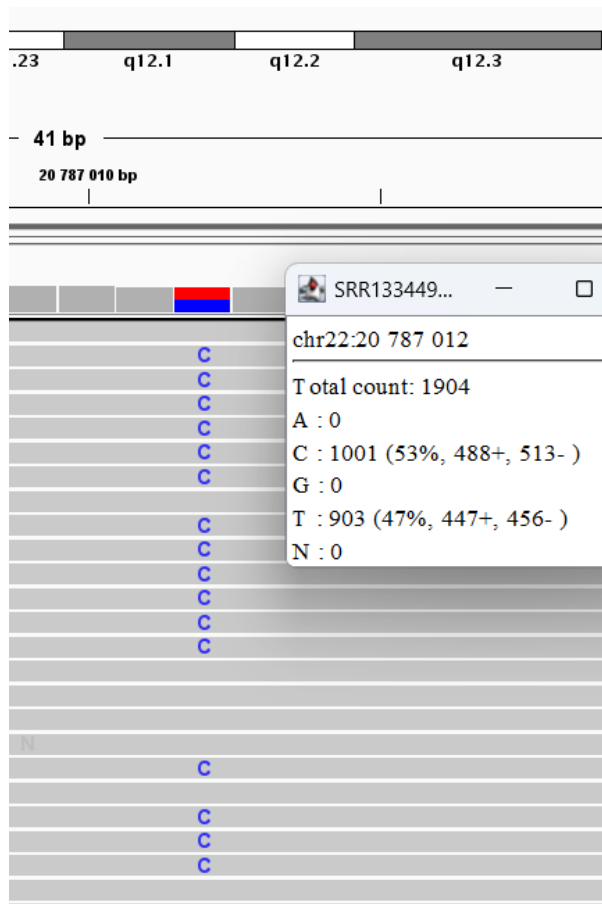
| Category | Illumina | | |
|----------|----------|--|---------|
| LR | | | Normal |
| LL | | | Unusual |
| RR | | | Unusual |
| RL | | | Unusual |

Pair orientation

<https://software.broadinstitute.org/software/igv/>

Visualization of Read Mappings

- Example of mutations, showing evidence



Practical example

- Show alignments sample_example
 - Position 9:21968200

Gather evidence for a mutation

- Number of reads, strand of reads, quality etc...

- Exercise:
 - Sample 1: 2:171868791
 - Sample 2: 6:121216049
 - Sample 3: X:129746638

Gather evidence for a mutation

- Number of reads, strand of reads, quality etc...